

02_end_to_end_machine_learning_project

February 21, 2023

Chapter 2 – End-to-end Machine Learning project

Welcome to Machine Learning Housing Corp.! Your task is to predict median house values in Californian districts, given a number of features from these districts.

This notebook contains all the sample code and solutions to the exercises in chapter 2.

Run in Google Colab

Warning: this is the code for the 1st edition of the book. Please visit <https://github.com/ageron/handson-ml2> for the 2nd edition code, with up-to-date notebooks using the latest library versions.

Note: You may find little differences between the code outputs in the book and in these Jupyter notebooks: these slight differences are mostly due to the random nature of many training algorithms: although I have tried to make these notebooks' outputs as constant as possible, it is impossible to guarantee that they will produce the exact same output on every platform. Also, some data structures (such as dictionaries) do not preserve the item order. Finally, I fixed a few minor bugs (I added notes next to the concerned cells) which lead to slightly different results, without changing the ideas presented in the book.

1 Setup

First, let's make sure this notebook works well in both python 2 and 3, import a few common modules, ensure Matplotlib plots figures inline and prepare a function to save the figures:

```
[1]: # To support both python 2 and python 3
    from __future__ import division, print_function, unicode_literals

    # Common imports
    import numpy as np
    import os

    # to make this notebook's output stable across runs
    np.random.seed(42)

    # To plot pretty figures
    %matplotlib inline
    import matplotlib as mpl
    import matplotlib.pyplot as plt
```

```

mpl.rcParams['axes', labelsizes=14)
mpl.rcParams['xtick', labelsizes=12)
mpl.rcParams['ytick', labelsizes=12)

# Where to save the figures
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "end_to_end_project"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

```

2 Get the data

```

[2]: import os
import tarfile
import urllib.request

DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml/master/"
HOUSING_PATH = os.path.join("datasets", "housing")
HOUSING_URL = DOWNLOAD_ROOT + "datasets/housing/housing.tgz"

def fetch_housing_data(housing_url=HOUSING_URL, housing_path=HOUSING_PATH):
    os.makedirs(housing_path, exist_ok=True)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(housing_url, tgz_path)
    housing_tgz = tarfile.open(tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()

```

```

[3]: fetch_housing_data()

```

```

[4]: import pandas as pd

def load_housing_data(housing_path=HOUSING_PATH):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)

```

```

[5]: housing = load_housing_data()
housing.head()

```

```
[5]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
[6]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[7]: housing["ocean_proximity"].value_counts()
```

```
[7]: <1H OCEAN      9136
INLAND          6551
NEAR OCEAN      2658
NEAR BAY        2290
ISLAND           5
Name: ocean_proximity, dtype: int64
```

```
[8]: housing.describe()
```

```
[8]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	

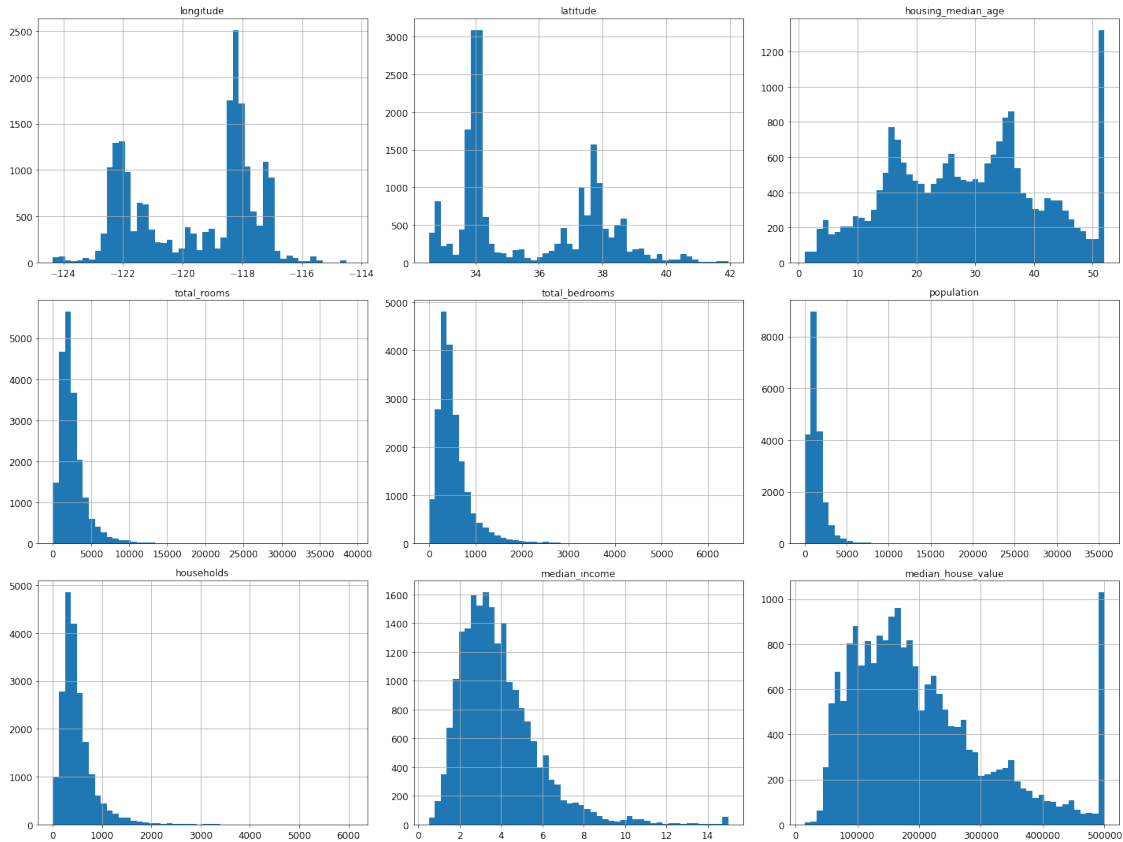
mean	-119.569704	35.631861	28.639486	2635.763081
std	2.003532	2.135952	12.585558	2181.615252
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[9]: %matplotlib inline
import matplotlib.pyplot as plt
housing.hist(bins=50, figsize=(20,15))
save_fig("attribute_histogram_plots")
plt.show()
```

Saving figure attribute_histogram_plots



```
[10]: # to make this notebook's output identical at every run
np.random.seed(42)
```

```
[11]: import numpy as np

# For illustration only. Sklearn has train_test_split()
def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

```
[12]: train_set, test_set = split_train_test(housing, 0.2)
print(len(train_set), "train +", len(test_set), "test")
```

16512 train + 4128 test

```
[13]: from zlib import crc32

def test_set_check(identifier, test_ratio):
```

```

        return crc32(np.int64(identifier)) & 0xffffffff < test_ratio * 2**32

def split_train_test_by_id(data, test_ratio, id_column):
    ids = data[id_column]
    in_test_set = ids.apply(lambda id_: test_set_check(id_, test_ratio))
    return data.loc[~in_test_set], data.loc[in_test_set]

```

The implementation of `test_set_check()` above works fine in both Python 2 and Python 3. In earlier releases, the following implementation was proposed, which supported any hash function, but was much slower and did not support Python 2:

```

[14]: import hashlib

def test_set_check(identifier, test_ratio, hash=hashlib.md5):
    return hash(np.int64(identifier)).digest()[-1] < 256 * test_ratio

```

If you want an implementation that supports any hash function and is compatible with both Python 2 and Python 3, here is one:

```

[15]: def test_set_check(identifier, test_ratio, hash=hashlib.md5):
        return bytearray(hash(np.int64(identifier)).digest())[-1] < 256 * test_ratio

```

```

[16]: housing_with_id = housing.reset_index()    # adds an `index` column
train_set, test_set = split_train_test_by_id(housing_with_id, 0.2, "index")

```

```

[17]: housing_with_id["id"] = housing["longitude"] * 1000 + housing["latitude"]
train_set, test_set = split_train_test_by_id(housing_with_id, 0.2, "id")

```

```

[18]: test_set.head()

```

```

[18]:      index  longitude  latitude  housing_median_age  total_rooms  \
8         8    -122.26    37.84             42.0         2555.0
10        10    -122.26    37.85             52.0         2202.0
11        11    -122.26    37.85             52.0         3503.0
12        12    -122.26    37.85             52.0         2491.0
13        13    -122.26    37.84             52.0          696.0

      total_bedrooms  population  households  median_income  median_house_value  \
8              665.0    1206.0      595.0         2.0804         226700.0
10             434.0     910.0      402.0         3.2031         281500.0
11             752.0    1504.0      734.0         3.2705         241800.0
12             474.0    1098.0      468.0         3.0750         213500.0
13             191.0     345.0      174.0         2.6736         191300.0

      ocean_proximity      id
8      NEAR BAY -122222.16
10     NEAR BAY -122222.15

```

```

11      NEAR BAY -122222.15
12      NEAR BAY -122222.15
13      NEAR BAY -122222.16

```

```

[19]: from sklearn.model_selection import train_test_split

      train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)

```

```

[20]: test_set.head()

```

```

[20]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
20046    -119.01    36.06             25.0         1505.0             NaN
3024     -119.46    35.14             30.0         2943.0             NaN
15663    -122.44    37.80             52.0         3830.0             NaN
20484    -118.72    34.28             17.0         3051.0             NaN
9814     -121.93    36.62             34.0         2351.0             NaN

```

```

      population  households  median_income  median_house_value  \
20046      1392.0       359.0         1.6812          47700.0
3024      1565.0       584.0         2.5313          45800.0
15663      1310.0       963.0         3.4801         500001.0
20484      1705.0       495.0         5.7376          218600.0
9814      1063.0       428.0         3.7250          278000.0

```

```

      ocean_proximity
20046          INLAND
3024          INLAND
15663          NEAR BAY
20484      <1H OCEAN
9814          NEAR OCEAN

```

```


[21]: housing["median_income"].hist()

```

```

[21]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3be9fcb730>

```



Warning: in the book, I did not use `pd.cut()`, instead I used the code below. The `pd.cut()` solution gives the same result (except the labels are integers instead of floats), but it is simpler to understand:

```
# Divide by 1.5 to limit the number of income categories
housing["income_cat"] = np.ceil(housing["median_income"] / 1.5)
# Label those above 5 as 5
housing["income_cat"].where(housing["income_cat"] < 5, 5.0, inplace=True)
```


```
[22]: housing["income_cat"] = pd.cut(housing["median_income"],
                                     bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                                     labels=[1, 2, 3, 4, 5])
```

```
[23]: housing["income_cat"].value_counts()
```

```
[23]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```

```
[24]: housing["income_cat"].hist()
```

```
[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3be9f289d0>
```

```
[25]: from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
```

```
[26]: strat_test_set["income_cat"].value_counts() / len(strat_test_set)
```

```
[26]: 3    0.350533
      2    0.318798
      4    0.176357
      5    0.114341
      1    0.039971
      Name: income_cat, dtype: float64
```

```
[27]: housing["income_cat"].value_counts() / len(housing)
```

```
[27]: 3    0.350581
      2    0.318847
      4    0.176308
      5    0.114438
      1    0.039826
      Name: income_cat, dtype: float64
```

```
[28]: def income_cat_proportions(data):
        return data["income_cat"].value_counts() / len(data)

train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)

compare_props = pd.DataFrame({
    "Overall": income_cat_proportions(housing),
    "Stratified": income_cat_proportions(strat_test_set),
    "Random": income_cat_proportions(test_set),
}).sort_index()
compare_props["Rand. %error"] = 100 * compare_props["Random"] /
    ↳compare_props["Overall"] - 100
compare_props["Strat. %error"] = 100 * compare_props["Stratified"] /
    ↳compare_props["Overall"] - 100
```

```
[29]: compare_props
```

```
[29]:
```

	Overall	Stratified	Random	Rand. %error	Strat. %error
1	0.039826	0.039971	0.040213	0.973236	0.364964
2	0.318847	0.318798	0.324370	1.732260	-0.015195
3	0.350581	0.350533	0.358527	2.266446	-0.013820
4	0.176308	0.176357	0.167393	-5.056334	0.027480
5	0.114438	0.114341	0.109496	-4.318374	-0.084674


```
[30]: for set_ in (strat_train_set, strat_test_set):
        set_.drop("income_cat", axis=1, inplace=True)
```

3 Discover and visualize the data to gain insights

```
[31]: housing = strat_train_set.copy()
```


```
[32]: housing.plot(kind="scatter", x="longitude", y="latitude")
save_fig("bad_visualization_plot")
```

Saving figure bad_visualization_plot



```
[33]: housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
      save_fig("better_visualization_plot")
```


Saving figure better_visualization_plot



The argument `sharex=False` fixes a display bug (the x-axis values and legend were not displayed). This is a temporary fix (see: <https://github.com/pandas-dev/pandas/issues/10611>). Thanks to Wilmer Arellano for pointing it out.

```
[34]: housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,  
    s=housing["population"]/100, label="population", figsize=(10,7),  
    c="median_house_value", cmap=plt.get_cmap("jet"), colorbar=True,  
    sharex=False)  
plt.legend()  
save_fig("housing_prices_scatterplot")
```

Saving figure `housing_prices_scatterplot`



```
[35]: # Download the California image
images_path = os.path.join(PROJECT_ROOT_DIR, "images", "end_to_end_project")
os.makedirs(images_path, exist_ok=True)
DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml/master/"
filename = "california.png"
print("Downloading", filename)
url = DOWNLOAD_ROOT + "images/end_to_end_project/" + filename
urllib.request.urlretrieve(url, os.path.join(images_path, filename))
```

Downloading california.png

```
[35]: ('./images/end_to_end_project/california.png',
<http.client.HTTPMessage at 0x7f3be9e50400>)
```

```
[36]: import matplotlib.image as mpimg
california_img=mpimg.imread(PROJECT_ROOT_DIR + '/images/end_to_end_project/
↳california.png')
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
s=housing['population']/100, label="Population",
c="median_house_value", cmap=plt.get_cmap("jet"),
colorbar=False, alpha=0.4,
)
```

```

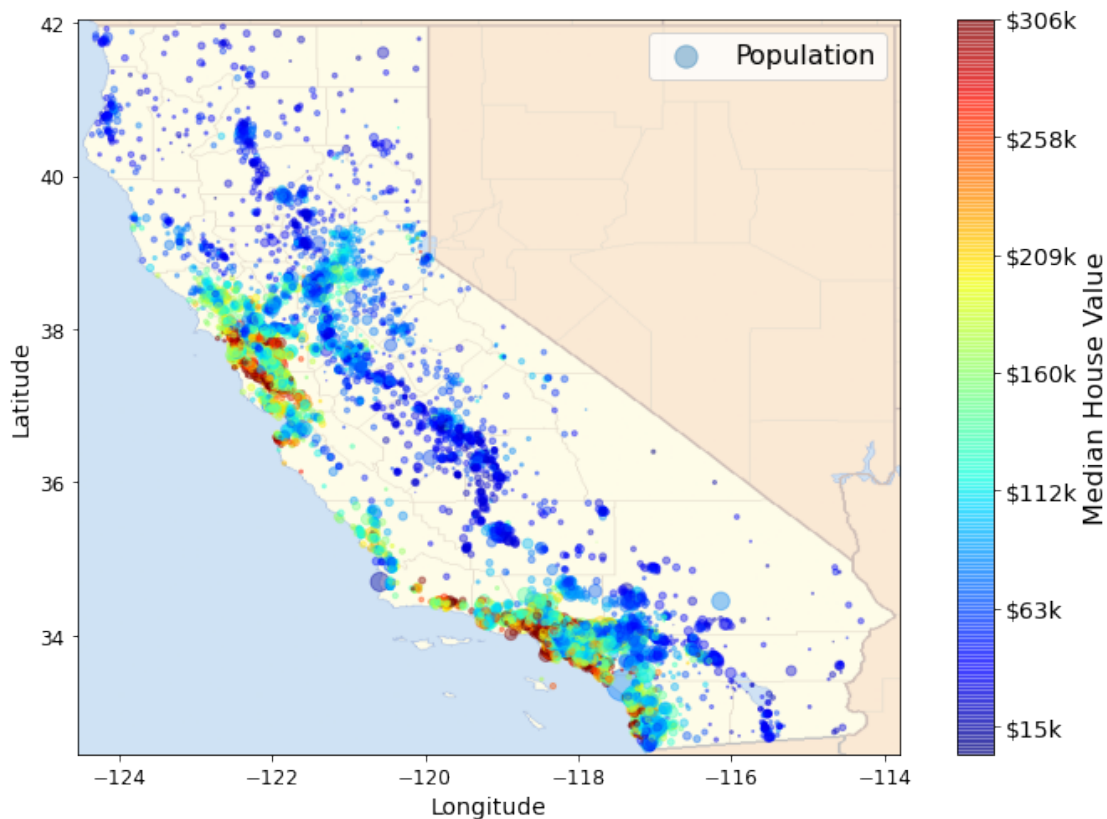
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cbar = plt.colorbar()
cbar.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
                        ↪ fontsize=14)
cbar.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()

```

Saving figure california_housing_prices_plot



```
[37]: corr_matrix = housing.corr()
```

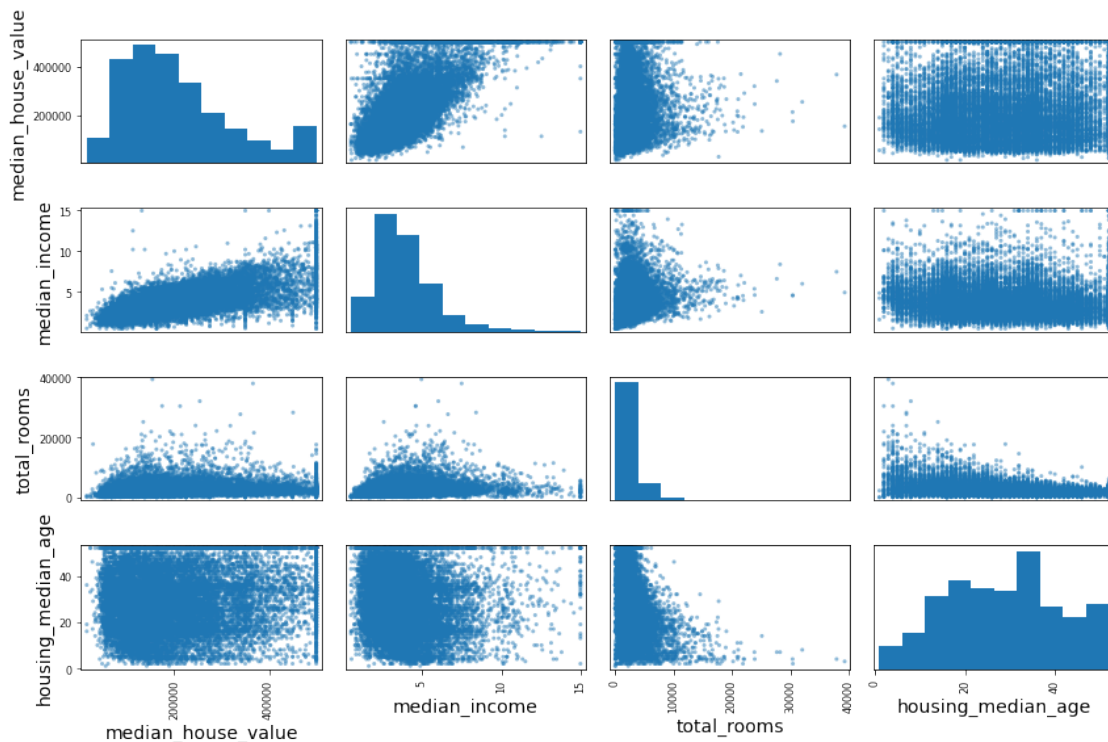
```
[38]: corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[38]: median_house_value    1.000000
      median_income        0.687151
      total_rooms          0.135140
      housing_median_age    0.114146
      households           0.064590
      total_bedrooms        0.047781
      population           -0.026882
      longitude            -0.047466
      latitude             -0.142673
      Name: median_house_value, dtype: float64
```

```
[39]: # from pandas.tools.plotting import scatter_matrix # For older versions of
      ↪Pandas
      from pandas.plotting import scatter_matrix


      attributes = ["median_house_value", "median_income", "total_rooms",
                    "housing_median_age"]
      scatter_matrix(housing[attributes], figsize=(12, 8))
      save_fig("scatter_matrix_plot")
```

Saving figure scatter_matrix_plot



```
[40]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                alpha=0.1)
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



```
[41]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

Note: there was a bug in the previous cell, in the definition of the `rooms_per_household` attribute. This explains why the correlation value below differs slightly from the value in the book (unless you are reading the latest version).

```
[42]: corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[42]: median_house_value    1.000000
median_income              0.687151
rooms_per_household        0.146255
total_rooms                0.135140
housing_median_age         0.114146
households                 0.064590
```



```


total_bedrooms          0.047781
population_per_household -0.021991
population               -0.026882
longitude               -0.047466
latitude                -0.142673
bedrooms_per_room       -0.259952
Name: median_house_value, dtype: float64

```

```

[43]: housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()

```



```

[44]: housing.describe()

```

```

[44]:
count    longitude    latitude    housing_median_age    total_rooms  \
count    16512.000000    16512.000000    16512.000000    16512.000000
mean      -119.575635     35.639314        28.653404     2622.539789
std         2.001828       2.137963        12.574819     2138.417080
min       -124.350000     32.540000         1.000000         6.000000
25%       -121.800000     33.940000        18.000000    1443.000000
50%       -118.510000     34.260000        29.000000    2119.000000
75%       -118.010000     37.720000        37.000000    3141.000000
max       -114.310000     41.950000        52.000000   39320.000000

```

	total_bedrooms	population	households	median_income \
count	16354.000000	16512.000000	16512.000000	16512.000000
mean	534.914639	1419.687379	497.011810	3.875884
std	412.665649	1115.663036	375.696156	1.904931
min	2.000000	3.000000	2.000000	0.499900
25%	295.000000	784.000000	279.000000	2.566950
50%	433.000000	1164.000000	408.000000	3.541550
75%	644.000000	1719.000000	602.000000	4.745325
max	6210.000000	35682.000000	5358.000000	15.000100

	median_house_value	rooms_per_household	bedrooms_per_room \
count	16512.000000	16512.000000	16354.000000
mean	207005.322372	5.440406	0.212873
std	115701.297250	2.611696	0.057378
min	14999.000000	1.130435	0.100000
25%	119800.000000	4.442168	0.175304
50%	179500.000000	5.232342	0.203027
75%	263900.000000	6.056361	0.239816
max	500001.000000	141.909091	1.000000

	population_per_household
count	16512.000000
mean	3.096469
std	11.584825
min	0.692308
25%	2.431352
50%	2.817661
75%	3.281420
max	1243.333333

4 Prepare the data for Machine Learning algorithms

```
[45]: housing = strat_train_set.drop("median_house_value", axis=1) # drop labels for
      ↪ training set
      housing_labels = strat_train_set["median_house_value"].copy()
```

```
[46]: sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
      sample_incomplete_rows
```

```
[46]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
1606	-122.08	37.88	26.0	2947.0	NaN	
10915	-117.87	33.73	45.0	2264.0	NaN	
19150	-122.70	38.35	14.0	2313.0	NaN	
4186	-118.23	34.13	48.0	1308.0	NaN	
16885	-122.40	37.58	26.0	3281.0	NaN	

	population	households	median_income	ocean_proximity
1606	825.0	626.0	2.9330	NEAR BAY
10915	1970.0	499.0	3.4193	<1H OCEAN
19150	954.0	397.0	3.7813	<1H OCEAN
4186	835.0	294.0	4.2891	<1H OCEAN
16885	1145.0	480.0	6.3580	NEAR OCEAN

```
[47]: sample_incomplete_rows.dropna(subset=["total_bedrooms"]) # option 1
```

```
[47]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
population, households, median_income, ocean_proximity]
Index: []
```

```
[48]: sample_incomplete_rows.drop("total_bedrooms", axis=1) # option 2
```

```
[48]:
```

	longitude	latitude	housing_median_age	total_rooms	population \
1606	-122.08	37.88	26.0	2947.0	825.0
10915	-117.87	33.73	45.0	2264.0	1970.0
19150	-122.70	38.35	14.0	2313.0	954.0
4186	-118.23	34.13	48.0	1308.0	835.0
16885	-122.40	37.58	26.0	3281.0	1145.0

	households	median_income	ocean_proximity
1606	626.0	2.9330	NEAR BAY
10915	499.0	3.4193	<1H OCEAN
19150	397.0	3.7813	<1H OCEAN
4186	294.0	4.2891	<1H OCEAN
16885	480.0	6.3580	NEAR OCEAN

```
[49]: median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option 3
sample_incomplete_rows
```

```
[49]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms \
1606	-122.08	37.88	26.0	2947.0	433.0
10915	-117.87	33.73	45.0	2264.0	433.0
19150	-122.70	38.35	14.0	2313.0	433.0
4186	-118.23	34.13	48.0	1308.0	433.0
16885	-122.40	37.58	26.0	3281.0	433.0

	population	households	median_income	ocean_proximity
1606	825.0	626.0	2.9330	NEAR BAY
10915	1970.0	499.0	3.4193	<1H OCEAN
19150	954.0	397.0	3.7813	<1H OCEAN
4186	835.0	294.0	4.2891	<1H OCEAN

```
16885      1145.0      480.0      6.3580      NEAR OCEAN
```

Warning: Since Scikit-Learn 0.20, the `sklearn.preprocessing.Imputer` class was replaced by the `sklearn.impute.SimpleImputer` class.

```
[50]: try:
      from sklearn.impute import SimpleImputer # Scikit-Learn 0.20+
    except ImportError:
      from sklearn.preprocessing import Imputer as SimpleImputer

    imputer = SimpleImputer(strategy="median")
```

Remove the text attribute because median can only be calculated on numerical attributes:

```
[51]: housing_num = housing.drop('ocean_proximity', axis=1)
      # alternatively: housing_num = housing.select_dtypes(include=[np.number])
```

```
[52]: imputer.fit(housing_num)
```

```
[52]: SimpleImputer(strategy='median')
```

```
[53]: imputer.statistics_
```

```
[53]: array([-118.51   ,   34.26   ,   29.        , 2119.        ,   433.        ,
          1164.        ,   408.        ,   3.54155])
```

Check that this is the same as manually computing the median of each attribute:

```
[54]: housing_num.median().values
```

```
[54]: array([-118.51   ,   34.26   ,   29.        , 2119.        ,   433.        ,
          1164.        ,   408.        ,   3.54155])
```

Transform the training set:

```
[55]: X = imputer.transform(housing_num)
```

```
[56]: housing_tr = pd.DataFrame(X, columns=housing_num.columns,
                              index=housing.index)
```

```
[57]: housing_tr.loc[sample_incomplete_rows.index.values]
```

```
[57]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
1606	-122.08	37.88	26.0	2947.0	433.0	
10915	-117.87	33.73	45.0	2264.0	433.0	
19150	-122.70	38.35	14.0	2313.0	433.0	
4186	-118.23	34.13	48.0	1308.0	433.0	
16885	-122.40	37.58	26.0	3281.0	433.0	

	population	households	median_income
1606	825.0	626.0	2.9330
10915	1970.0	499.0	3.4193
19150	954.0	397.0	3.7813
4186	835.0	294.0	4.2891
16885	1145.0	480.0	6.3580

```
[58]: imputer.strategy
```

```
[58]: 'median'
```

```
[59]: housing_tr = pd.DataFrame(X, columns=housing_num.columns,
                                index=housing_num.index)
housing_tr.head()
```

```
[59]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
12655    -121.46    38.52           29.0        3873.0         797.0
15502    -117.23    33.09           7.0        5320.0         855.0
2908     -119.04    35.37          44.0        1618.0         310.0
14053    -117.13    32.75          24.0        1877.0         519.0
20496    -118.70    34.28          27.0        3536.0         646.0
```

	population	households	median_income
12655	2237.0	706.0	2.1736
15502	2015.0	768.0	6.3373
2908	667.0	300.0	2.8750
14053	898.0	483.0	2.2264
20496	1837.0	580.0	4.4964

Now let's preprocess the categorical input feature, `ocean_proximity`:

```
[60]: housing_cat = housing[['ocean_proximity']]
housing_cat.head(10)
```

```
[60]:      ocean_proximity
12655          INLAND
15502      NEAR OCEAN
2908          INLAND
14053      NEAR OCEAN
20496    <1H OCEAN
1481          NEAR BAY
18125    <1H OCEAN
5830      <1H OCEAN
17989    <1H OCEAN
4861      <1H OCEAN
```

Warning: earlier versions of the book used the `LabelEncoder` class or `Pandas`'

`Series.factorize()` method to encode string categorical attributes as integers. However, the `OrdinalEncoder` class that was introduced in Scikit-Learn 0.20 (see [PR #10521](#)) is preferable since it is designed for input features (X instead of labels y) and it plays well with pipelines (introduced later in this notebook). If you are using an older version of Scikit-Learn (<0.20), then you can import it from `future_encoders.py` instead.

```
[61]: try:
        from sklearn.preprocessing import OrdinalEncoder
    except ImportError:
        from future_encoders import OrdinalEncoder # Scikit-Learn < 0.20
```

```
[62]: ordinal_encoder = OrdinalEncoder()
housing_cat_encoded = ordinal_encoder.fit_transform(housing_cat)
housing_cat_encoded[:10]
```

```
[62]: array([[1.],
           [4.],
           [1.],
           [4.],
           [0.],
           [3.],
           [0.],
           [0.],
           [0.],
           [0.]])
```

```
[63]: ordinal_encoder.categories_
```

```
[63]: [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
           dtype=object)]
```

Warning: earlier versions of the book used the `LabelBinarizer` or `CategoricalEncoder` classes to convert each categorical value to a one-hot vector. It is now preferable to use the `OneHotEncoder` class. Since Scikit-Learn 0.20 it can handle string categorical inputs (see [PR #10521](#)), not just integer categorical inputs. If you are using an older version of Scikit-Learn, you can import the new version from `future_encoders.py`:

```
[64]: try:
        from sklearn.preprocessing import OrdinalEncoder # just to raise an
        ↳ ImportError if Scikit-Learn < 0.20
        from sklearn.preprocessing import OneHotEncoder
    except ImportError:
        from future_encoders import OneHotEncoder # Scikit-Learn < 0.20

    cat_encoder = OneHotEncoder()
    housing_cat_1hot = cat_encoder.fit_transform(housing_cat)
    housing_cat_1hot
```

```
[64]: <16512x5 sparse matrix of type '<class 'numpy.float64'>'
      with 16512 stored elements in Compressed Sparse Row format>
```

By default, the `OneHotEncoder` class returns a sparse array, but we can convert it to a dense array if needed by calling the `toarray()` method:

```
[65]: housing_cat_1hot.toarray()
```

```
[65]: array([[0., 1., 0., 0., 0.],
           [0., 0., 0., 0., 1.],
           [0., 1., 0., 0., 0.],
           ...,
           [1., 0., 0., 0., 0.],
           [1., 0., 0., 0., 0.],
           [0., 1., 0., 0., 0.]])
```

Alternatively, you can set `sparse=False` when creating the `OneHotEncoder`:

```
[66]: cat_encoder = OneHotEncoder(sparse=False)
      housing_cat_1hot = cat_encoder.fit_transform(housing_cat)
      housing_cat_1hot
```

```
[66]: array([[0., 1., 0., 0., 0.],
           [0., 0., 0., 0., 1.],
           [0., 1., 0., 0., 0.],
           ...,
           [1., 0., 0., 0., 0.],
           [1., 0., 0., 0., 0.],
           [0., 1., 0., 0., 0.]])
```

```
[67]: cat_encoder.categories_
```

```
[67]: [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
           dtype=object)]
```

Let's create a custom transformer to add extra attributes:

```
[68]: housing.columns
```

```
[68]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
           'total_bedrooms', 'population', 'households', 'median_income',
           'ocean_proximity'],
          dtype='object')
```

```
[69]: from sklearn.base import BaseEstimator, TransformerMixin

      # get the right column indices: safer than hard-coding indices 3, 4, 5, 6
      rooms_ix, bedrooms_ix, population_ix, household_ix = [
```

```

list(housing.columns).index(col)
for col in ("total_rooms", "total_bedrooms", "population", "households")]

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
    def __init__(self, add_bedrooms_per_room = True): # no *args or **kwargs
        self.add_bedrooms_per_room = add_bedrooms_per_room
    def fit(self, X, y=None):
        return self # nothing else to do
    def transform(self, X, y=None):
        rooms_per_household = X[:, rooms_ix] / X[:, household_ix]
        population_per_household = X[:, population_ix] / X[:, household_ix]
        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
            return np.c_[X, rooms_per_household, population_per_household,
                          bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]

attr_adder = CombinedAttributesAdder(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values)

```

Alternatively, you can use Scikit-Learn's `FunctionTransformer` class that lets you easily create a transformer based on a transformation function (thanks to [Hanmin Qin](#) for suggesting this code). Note that we need to set `validate=False` because the data contains non-float values (`validate` will default to `False` in Scikit-Learn 0.22).

```

[70]: from sklearn.preprocessing import FunctionTransformer

def add_extra_features(X, add_bedrooms_per_room=True):
    rooms_per_household = X[:, rooms_ix] / X[:, household_ix]
    population_per_household = X[:, population_ix] / X[:, household_ix]
    if add_bedrooms_per_room:
        bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
        return np.c_[X, rooms_per_household, population_per_household,
                      bedrooms_per_room]
    else:
        return np.c_[X, rooms_per_household, population_per_household]

attr_adder = FunctionTransformer(add_extra_features, validate=False,
                                 kw_args={"add_bedrooms_per_room": False})
housing_extra_attribs = attr_adder.fit_transform(housing.values)

```

```

[71]: housing_extra_attribs = pd.DataFrame(
    housing_extra_attribs,
    columns=list(housing.columns)+["rooms_per_household",
    ↪ "population_per_household"],
    index=housing.index)

```



```
housing_extra_attribs.head()
```

```
[71]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
12655	-121.46	38.52	29.0	3873.0	797.0	
15502	-117.23	33.09	7.0	5320.0	855.0	
2908	-119.04	35.37	44.0	1618.0	310.0	
14053	-117.13	32.75	24.0	1877.0	519.0	
20496	-118.7	34.28	27.0	3536.0	646.0	

	population	households	median_income	ocean_proximity	rooms_per_household	\
12655	2237.0	706.0	2.1736	INLAND	5.485836	
15502	2015.0	768.0	6.3373	NEAR OCEAN	6.927083	
2908	667.0	300.0	2.875	INLAND	5.393333	
14053	898.0	483.0	2.2264	NEAR OCEAN	3.886128	
20496	1837.0	580.0	4.4964	<1H OCEAN	6.096552	

	population_per_household
12655	3.168555
15502	2.623698
2908	2.223333
14053	1.859213
20496	3.167241

Now let's build a pipeline for preprocessing the numerical attributes (note that we could use `CombinedAttributesAdder()` instead of `FunctionTransformer(...)` if we preferred):

```
[72]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler

      num_pipeline = Pipeline([
          ('imputer', SimpleImputer(strategy="median")),
          ('attribs_adder', FunctionTransformer(add_extra_features,
      ↪ validate=False)),
          ('std_scaler', StandardScaler()),
      ])

      housing_num_tr = num_pipeline.fit_transform(housing_num)
```

```
[73]: housing_num_tr
```

```
[73]: array([[ -0.94135046,  1.34743822,  0.02756357, ...,  0.01739526,
          0.00622264, -0.12112176],
          [ 1.17178212, -1.19243966, -1.72201763, ...,  0.56925554,
          -0.04081077, -0.81086696],
          [ 0.26758118, -0.1259716 ,  1.22045984, ..., -0.01802432,
          -0.07537122, -0.33827252],
          ...,
```

```

[-1.5707942 ,  1.31001828,  1.53856552, ..., -0.5092404 ,
 -0.03743619,  0.32286937],
[-1.56080303,  1.2492109 , -1.1653327 , ...,  0.32814891,
 -0.05915604, -0.45702273],
[-1.28105026,  2.02567448, -0.13148926, ...,  0.01407228,
 0.00657083, -0.12169672]])

```

Warning: earlier versions of the book applied different transformations to different columns using a solution based on a `DataFrameSelector` transformer and a `FeatureUnion` (see below). It is now preferable to use the `ColumnTransformer` class that was introduced in Scikit-Learn 0.20. If you are using an older version of Scikit-Learn, you can import it from `future_encoders.py`:

```

[74]: try:
        from sklearn.compose import ColumnTransformer
    except ImportError:
        from future_encoders import ColumnTransformer # Scikit-Learn < 0.20

```

```

[75]: num_attribs = list(housing_num)
    cat_attribs = ["ocean_proximity"]

    full_pipeline = ColumnTransformer([
        ("num", num_pipeline, num_attribs),
        ("cat", OneHotEncoder(), cat_attribs),
    ])

    housing_prepared = full_pipeline.fit_transform(housing)

```

```

[76]: housing_prepared

```

```

[76]: array([[ -0.94135046,  1.34743822,  0.02756357, ...,  0.          ,
         0.          ,  0.          ],
 [ 1.17178212, -1.19243966, -1.72201763, ...,  0.          ,
         0.          ,  1.          ],
 [ 0.26758118, -0.1259716 ,  1.22045984, ...,  0.          ,
         0.          ,  0.          ],
 ...,
 [-1.5707942 ,  1.31001828,  1.53856552, ...,  0.          ,
         0.          ,  0.          ],
 [-1.56080303,  1.2492109 , -1.1653327 , ...,  0.          ,
         0.          ,  0.          ],
 [-1.28105026,  2.02567448, -0.13148926, ...,  0.          ,
         0.          ,  0.          ]])

```

```

[77]: housing_prepared.shape

```

```

[77]: (16512, 16)

```

For reference, here is the old solution based on a `DataFrameSelector` transformer (to just select a

subset of the Pandas DataFrame columns), and a FeatureUnion:

```
[78]: from sklearn.base import BaseEstimator, TransformerMixin

# Create a class to select numerical or categorical columns
class OldDataFrameSelector(BaseEstimator, TransformerMixin):
    def __init__(self, attribute_names):
        self.attribute_names = attribute_names
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        return X[self.attribute_names].values
```

Now let's join all these components into a big pipeline that will preprocess both the numerical and the categorical features (again, we could use `CombinedAttributesAdder()` instead of `FunctionTransformer(...)` if we preferred):

```
[79]: num_attribs = list(housing_num)
cat_attribs = ["ocean_proximity"]

old_num_pipeline = Pipeline([
    ('selector', OldDataFrameSelector(num_attribs)),
    ('imputer', SimpleImputer(strategy="median")),
    ('attrs_adder', FunctionTransformer(add_extra_features,
    ↪ validate=False)),
    ('std_scaler', StandardScaler()),
])

old_cat_pipeline = Pipeline([
    ('selector', OldDataFrameSelector(cat_attribs)),
    ('cat_encoder', OneHotEncoder(sparse=False)),
])
```

```
[80]: from sklearn.pipeline import FeatureUnion

old_full_pipeline = FeatureUnion(transformer_list=[
    ("num_pipeline", old_num_pipeline),
    ("cat_pipeline", old_cat_pipeline),
])
```

```
[81]: old_housing_prepared = old_full_pipeline.fit_transform(housing)
old_housing_prepared
```

```
[81]: array([[ -0.94135046,  1.34743822,  0.02756357, ...,  0.          ,
                0.          ,  0.          ],
       [ 1.17178212, -1.19243966, -1.72201763, ...,  0.          ,
                0.          ,  1.          ],
       [ 0.26758118, -0.1259716 ,  1.22045984, ...,  0.          ,
```

```

0.          , 0.          ],
...,
[-1.5707942 , 1.31001828, 1.53856552, ..., 0.          ,
 0.          , 0.          ],
[-1.56080303, 1.2492109 , -1.1653327 , ..., 0.          ,
 0.          , 0.          ],
[-1.28105026, 2.02567448, -0.13148926, ..., 0.          ,
 0.          , 0.          ]])

```

The result is the same as with the ColumnTransformer:

```
[82]: np.allclose(housing_prepared, old_housing_prepared)
```

```
[82]: True
```

5 Select and train a model

```
[83]: from sklearn.linear_model import LinearRegression
```

```

lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)

```

```
[83]: LinearRegression()
```

```

[84]: # let's try the full preprocessing pipeline on a few training instances
some_data = housing.iloc[:5]
some_labels = housing_labels.iloc[:5]
some_data_prepared = full_pipeline.transform(some_data)

print("Predictions:", lin_reg.predict(some_data_prepared))

```

```
Predictions: [ 85657.90192014 305492.60737488 152056.46122456 186095.70946094
244550.67966089]
```

Compare against the actual values:

```
[85]: print("Labels:", list(some_labels))
```

```
Labels: [72100.0, 279600.0, 82700.0, 112500.0, 238300.0]
```

```
[86]: some_data_prepared
```

```

[86]: array([[ -0.94135046,  1.34743822,  0.02756357,  0.58477745,  0.64037127,
         0.73260236,  0.55628602, -0.8936472 ,  0.01739526,  0.00622264,
        -0.12112176,  0.          ,  1.          ,  0.          ,  0.          ,
         0.          ],
        [ 1.17178212, -1.19243966, -1.72201763,  1.26146668,  0.78156132,
         0.53361152,  0.72131799,  1.292168 ,  0.56925554, -0.04081077,

```

```

-0.81086696, 0.          , 0.          , 0.          , 0.          ,
1.          ],
[ 0.26758118, -0.1259716 , 1.22045984, -0.46977281, -0.54513828,
-0.67467519, -0.52440722, -0.52543365, -0.01802432, -0.07537122,
-0.33827252, 0.          , 1.          , 0.          , 0.          ,
0.          ],
[ 1.22173797, -1.35147437, -0.37006852, -0.34865152, -0.03636724,
-0.46761716, -0.03729672, -0.86592882, -0.59513997, -0.10680295,
0.96120521, 0.          , 0.          , 0.          , 0.          ,
1.          ],
[ 0.43743108, -0.63581817, -0.13148926, 0.42717947, 0.27279028,
0.37406031, 0.22089846, 0.32575178, 0.2512412 , 0.00610923,
-0.47451338, 1.          , 0.          , 0.          , 0.          ,
0.          ]]

```

```
[87]: from sklearn.metrics import mean_squared_error
```

```

housing_predictions = lin_reg.predict(housing_prepared)
lin_mse = mean_squared_error(housing_labels, housing_predictions)
lin_rmse = np.sqrt(lin_mse)
lin_rmse

```

```
[87]: 68627.87390018745
```

```
[88]: from sklearn.metrics import mean_absolute_error
```

```

lin_mae = mean_absolute_error(housing_labels, housing_predictions)
lin_mae

```

```
[88]: 49438.66860915802
```

```
[89]: from sklearn.tree import DecisionTreeRegressor
```

```

tree_reg = DecisionTreeRegressor(random_state=42)
tree_reg.fit(housing_prepared, housing_labels)

```

```
[89]: DecisionTreeRegressor(random_state=42)
```

```

[90]: housing_predictions = tree_reg.predict(housing_prepared)
tree_mse = mean_squared_error(housing_labels, housing_predictions)
tree_rmse = np.sqrt(tree_mse)
tree_rmse

```

```
[90]: 0.0
```

6 Fine-tune your model

```
[91]: from sklearn.model_selection import cross_val_score

scores = cross_val_score(tree_reg, housing_prepared, housing_labels,
                          scoring="neg_mean_squared_error", cv=10)
tree_rmse_scores = np.sqrt(-scores)
```

```
[92]: def display_scores(scores):
        print("Scores:", scores)
        print("Mean:", scores.mean())
        print("Standard deviation:", scores.std())

display_scores(tree_rmse_scores)
```

```
Scores: [72831.45749112 69973.18438322 69528.56551415 72517.78229792
 69145.50006909 79094.74123727 68960.045444 73344.50225684
 69826.02473916 71077.09753998]
Mean: 71629.89009727491
Standard deviation: 2914.035468468928
```

```
[93]: lin_scores = cross_val_score(lin_reg, housing_prepared, housing_labels,
                                   scoring="neg_mean_squared_error", cv=10)
lin_rmse_scores = np.sqrt(-lin_scores)
display_scores(lin_rmse_scores)
```

```
Scores: [71762.76364394 64114.99166359 67771.17124356 68635.19072082
 66846.14089488 72528.03725385 73997.08050233 68802.33629334
 66443.28836884 70139.79923956]
Mean: 69104.07998247063
Standard deviation: 2880.3282098180634
```

Note: we specify `n_estimators=10` to avoid a warning about the fact that the default value is going to change to 100 in Scikit-Learn 0.22.

```
[94]: from sklearn.ensemble import RandomForestRegressor

forest_reg = RandomForestRegressor(n_estimators=10, random_state=42)
forest_reg.fit(housing_prepared, housing_labels)
```

```
[94]: RandomForestRegressor(n_estimators=10, random_state=42)
```

```
[95]: housing_predictions = forest_reg.predict(housing_prepared)
forest_mse = mean_squared_error(housing_labels, housing_predictions)
forest_rmse = np.sqrt(forest_mse)
forest_rmse
```

```
[95]: 22413.454658589766
```

```
[96]: from sklearn.model_selection import cross_val_score

forest_scores = cross_val_score(forest_reg, housing_prepared, housing_labels,
                                scoring="neg_mean_squared_error", cv=10)
forest_rmse_scores = np.sqrt(-forest_scores)
display_scores(forest_rmse_scores)
```

```
Scores: [53519.05518628 50467.33817051 48924.16513902 53771.72056856
50810.90996358 54876.09682033 56012.79985518 52256.88927227
51527.73185039 55762.56008531]
Mean: 52792.92669114079
Standard deviation: 2262.8151900582
```

```
[97]: scores = cross_val_score(lin_reg, housing_prepared, housing_labels,
    ↪scoring="neg_mean_squared_error", cv=10)
pd.Series(np.sqrt(-scores)).describe()
```

```
[97]: count      10.000000
mean      69104.079982
std       3036.132517
min       64114.991664
25%      67077.398482
50%      68718.763507
75%      71357.022543
max       73997.080502
dtype: float64
```

```
[98]: from sklearn.svm import SVR

svm_reg = SVR(kernel="linear")
svm_reg.fit(housing_prepared, housing_labels)
housing_predictions = svm_reg.predict(housing_prepared)
svm_mse = mean_squared_error(housing_labels, housing_predictions)
svm_rmse = np.sqrt(svm_mse)
svm_rmse
```

```
[98]: 111095.06635291968
```

```
[99]: from sklearn.model_selection import GridSearchCV

param_grid = [
    # try 12 (3×4) combinations of hyperparameters
    {'n_estimators': [3, 10, 30], 'max_features': [2, 4, 6, 8]},
    # then try 6 (2×3) combinations with bootstrap set as False
    {'bootstrap': [False], 'n_estimators': [3, 10], 'max_features': [2, 3, 4]},
]
```

```

forest_reg = RandomForestRegressor(random_state=42)
# train across 5 folds, that's a total of (12+6)*5=90 rounds of training
grid_search = GridSearchCV(forest_reg, param_grid, cv=5,
                           scoring='neg_mean_squared_error',
                           ↪return_train_score=True)
grid_search.fit(housing_prepared, housing_labels)

```

```

[99]: GridSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
                param_grid=[{'max_features': [2, 4, 6, 8],
                              'n_estimators': [3, 10, 30]},
                              {'bootstrap': [False], 'max_features': [2, 3, 4],
                               'n_estimators': [3, 10]}],
                return_train_score=True, scoring='neg_mean_squared_error')

```

The best hyperparameter combination found:

```

[100]: grid_search.best_params_

```

```

[100]: {'max_features': 8, 'n_estimators': 30}

```

```

[101]: grid_search.best_estimator_

```

```

[101]: RandomForestRegressor(max_features=8, n_estimators=30, random_state=42)

```

Let's look at the score of each hyperparameter combination tested during the grid search:

```

[102]: cvres = grid_search.cv_results_
        for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
            print(np.sqrt(-mean_score), params)

```

```

63895.161577951665 {'max_features': 2, 'n_estimators': 3}
54916.32386349543 {'max_features': 2, 'n_estimators': 10}
52885.86715332332 {'max_features': 2, 'n_estimators': 30}
60075.3680329983 {'max_features': 4, 'n_estimators': 3}
52495.01284985185 {'max_features': 4, 'n_estimators': 10}
50187.24324926565 {'max_features': 4, 'n_estimators': 30}
58064.73529982314 {'max_features': 6, 'n_estimators': 3}
51519.32062366315 {'max_features': 6, 'n_estimators': 10}
49969.80441627874 {'max_features': 6, 'n_estimators': 30}
58895.824998155826 {'max_features': 8, 'n_estimators': 3}
52459.79624724529 {'max_features': 8, 'n_estimators': 10}
49898.98913455217 {'max_features': 8, 'n_estimators': 30}
62381.765106921855 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}
54476.57050944266 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}
59974.60028085155 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}
52754.5632813202 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}
57831.136061214274 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}
51278.37877140253 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}

```



```
[103]: pd.DataFrame(grid_search.cv_results_)
```

```
[103]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	0.072480	0.002602	0.004944	0.000461	
1	0.229180	0.004078	0.013138	0.000674	
2	0.672476	0.001871	0.035402	0.000784	
3	0.128708	0.021488	0.005058	0.000651	
4	0.518286	0.063770	0.014502	0.000863	
5	1.108341	0.002866	0.036428	0.002273	
6	0.150071	0.001643	0.005004	0.000607	
7	0.526446	0.036520	0.013911	0.001839	
8	1.700406	0.299065	0.037289	0.003874	
9	0.200003	0.003256	0.004759	0.000206	
10	0.794997	0.154712	0.014967	0.002049	
11	2.572026	0.575364	0.056408	0.038263	
12	0.109426	0.005002	0.005440	0.000106	
13	0.355580	0.006526	0.015074	0.000322	
14	0.141472	0.001391	0.005395	0.000078	
15	0.472482	0.005227	0.014986	0.000172	
16	0.181207	0.003012	0.005575	0.000345	
17	0.690830	0.128061	0.016427	0.001862	

	param_max_features	param_n_estimators	param_bootstrap	\
0	2	3	NaN	
1	2	10	NaN	
2	2	30	NaN	
3	4	3	NaN	
4	4	10	NaN	
5	4	30	NaN	
6	6	3	NaN	
7	6	10	NaN	
8	6	30	NaN	
9	8	3	NaN	
10	8	10	NaN	
11	8	30	NaN	
12	2	3	False	
13	2	10	False	
14	3	3	False	
15	3	10	False	
16	4	3	False	
17	4	10	False	

	params	split0_test_score	\
0	{'max_features': 2, 'n_estimators': 3}	-4.119912e+09	
1	{'max_features': 2, 'n_estimators': 10}	-2.973521e+09	
2	{'max_features': 2, 'n_estimators': 30}	-2.801229e+09	
3	{'max_features': 4, 'n_estimators': 3}	-3.528743e+09	

4	{'max_features': 4, 'n_estimators': 10}	-2.742620e+09
5	{'max_features': 4, 'n_estimators': 30}	-2.522176e+09
6	{'max_features': 6, 'n_estimators': 3}	-3.362127e+09
7	{'max_features': 6, 'n_estimators': 10}	-2.622099e+09
8	{'max_features': 6, 'n_estimators': 30}	-2.446142e+09
9	{'max_features': 8, 'n_estimators': 3}	-3.590333e+09
10	{'max_features': 8, 'n_estimators': 10}	-2.721311e+09
11	{'max_features': 8, 'n_estimators': 30}	-2.492636e+09
12	{'bootstrap': False, 'max_features': 2, 'n_est...	-4.020842e+09
13	{'bootstrap': False, 'max_features': 2, 'n_est...	-2.901352e+09
14	{'bootstrap': False, 'max_features': 3, 'n_est...	-3.687132e+09
15	{'bootstrap': False, 'max_features': 3, 'n_est...	-2.837028e+09
16	{'bootstrap': False, 'max_features': 4, 'n_est...	-3.549428e+09
17	{'bootstrap': False, 'max_features': 4, 'n_est...	-2.692499e+09

	split1_test_score	...	mean_test_score	std_test_score	rank_test_score	\
0	-3.723465e+09	...	-4.082592e+09	1.867375e+08	18	
1	-2.810319e+09	...	-3.015803e+09	1.139808e+08	11	
2	-2.671474e+09	...	-2.796915e+09	7.980892e+07	9	
3	-3.490303e+09	...	-3.609050e+09	1.375683e+08	16	
4	-2.609311e+09	...	-2.755726e+09	1.182604e+08	7	
5	-2.440241e+09	...	-2.518759e+09	8.488084e+07	3	
6	-3.311863e+09	...	-3.371513e+09	1.378086e+08	13	
7	-2.669655e+09	...	-2.654240e+09	6.967978e+07	5	
8	-2.446594e+09	...	-2.496981e+09	7.357046e+07	2	
9	-3.232664e+09	...	-3.468718e+09	1.293758e+08	14	
10	-2.675886e+09	...	-2.752030e+09	6.258030e+07	6	
11	-2.444818e+09	...	-2.489909e+09	7.086483e+07	1	
12	-3.951861e+09	...	-3.891485e+09	8.648595e+07	17	
13	-3.036875e+09	...	-2.967697e+09	4.582448e+07	10	
14	-3.446245e+09	...	-3.596953e+09	8.011960e+07	15	
15	-2.619558e+09	...	-2.783044e+09	8.862580e+07	8	
16	-3.318176e+09	...	-3.344440e+09	1.099355e+08	12	
17	-2.542704e+09	...	-2.629472e+09	8.510266e+07	4	

	split0_train_score	split1_train_score	split2_train_score	\
0	-1.155630e+09	-1.089726e+09	-1.153843e+09	
1	-5.982947e+08	-5.904781e+08	-6.123850e+08	
2	-4.412567e+08	-4.326398e+08	-4.553722e+08	
3	-9.782368e+08	-9.806455e+08	-1.003780e+09	
4	-5.063215e+08	-5.257983e+08	-5.081984e+08	
5	-3.776568e+08	-3.902106e+08	-3.885042e+08	
6	-8.909397e+08	-9.583733e+08	-9.000201e+08	
7	-4.939906e+08	-5.145996e+08	-5.023512e+08	
8	-3.760968e+08	-3.876636e+08	-3.875307e+08	
9	-9.505012e+08	-9.166119e+08	-9.033910e+08	
10	-4.998373e+08	-4.997970e+08	-5.099880e+08	

11	-3.801679e+08	-3.832972e+08	-3.823818e+08
12	-0.000000e+00	-4.306828e+01	-1.051392e+04
13	-0.000000e+00	-3.876145e+00	-9.462528e+02
14	-0.000000e+00	-0.000000e+00	-0.000000e+00
15	-0.000000e+00	-0.000000e+00	-0.000000e+00
16	-0.000000e+00	-0.000000e+00	-0.000000e+00
17	-0.000000e+00	-0.000000e+00	-0.000000e+00

	split3_train_score	split4_train_score	mean_train_score	std_train_score
0	-1.118149e+09	-1.093446e+09	-1.122159e+09	2.834288e+07
1	-5.727681e+08	-5.905210e+08	-5.928894e+08	1.284978e+07
2	-4.320746e+08	-4.311606e+08	-4.385008e+08	9.184397e+06
3	-1.016515e+09	-1.011270e+09	-9.980896e+08	1.577372e+07
4	-5.174405e+08	-5.282066e+08	-5.171931e+08	8.882622e+06
5	-3.830866e+08	-3.894779e+08	-3.857872e+08	4.774229e+06
6	-8.964731e+08	-9.151927e+08	-9.121998e+08	2.444837e+07
7	-4.959467e+08	-5.147087e+08	-5.043194e+08	8.880106e+06
8	-3.760938e+08	-3.861056e+08	-3.826981e+08	5.418747e+06
9	-9.070642e+08	-9.459386e+08	-9.247014e+08	1.973471e+07
10	-5.047868e+08	-5.348043e+08	-5.098427e+08	1.303601e+07
11	-3.778452e+08	-3.817589e+08	-3.810902e+08	1.916605e+06
12	-0.000000e+00	-0.000000e+00	-2.111398e+03	4.201294e+03
13	-0.000000e+00	-0.000000e+00	-1.900258e+02	3.781165e+02
14	-0.000000e+00	-0.000000e+00	0.000000e+00	0.000000e+00
15	-0.000000e+00	-0.000000e+00	0.000000e+00	0.000000e+00
16	-0.000000e+00	-0.000000e+00	0.000000e+00	0.000000e+00
17	-0.000000e+00	-0.000000e+00	0.000000e+00	0.000000e+00

[18 rows x 23 columns]

```
[104]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_distributions = {
    'n_estimators': randint(low=1, high=200),
    'max_features': randint(low=1, high=8),
}

forest_reg = RandomForestRegressor(random_state=42)
rnd_search = RandomizedSearchCV(forest_reg, param_distributions=param_distributions,
                                n_iter=10, cv=5,
                                scoring='neg_mean_squared_error', random_state=42)
rnd_search.fit(housing_prepared, housing_labels)
```

```
[104]: RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
                          param_distributions={'max_features':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f3be989d4c0>,
```

```

        'n_estimators':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f3be99fa490>},
        random_state=42, scoring='neg_mean_squared_error')

```

```

[105]: cvres = rnd_search.cv_results_
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
    print(np.sqrt(-mean_score), params)

```

```

49117.55344336652 {'max_features': 7, 'n_estimators': 180}
51450.63202856348 {'max_features': 5, 'n_estimators': 15}
50692.53588182537 {'max_features': 3, 'n_estimators': 72}
50783.614493515 {'max_features': 5, 'n_estimators': 21}
49162.89877456354 {'max_features': 7, 'n_estimators': 122}
50655.798471042704 {'max_features': 3, 'n_estimators': 75}
50513.856319990606 {'max_features': 3, 'n_estimators': 88}
49521.17201976928 {'max_features': 5, 'n_estimators': 100}
50302.90440763418 {'max_features': 3, 'n_estimators': 150}
65167.02018649492 {'max_features': 5, 'n_estimators': 2}

```

```

[106]: feature_importances = grid_search.best_estimator_.feature_importances_
feature_importances

```

```

[106]: array([6.96542523e-02, 6.04213840e-02, 4.21882202e-02, 1.52450557e-02,
1.55545295e-02, 1.58491147e-02, 1.49346552e-02, 3.79009225e-01,
5.47789150e-02, 1.07031322e-01, 4.82031213e-02, 6.79266007e-03,
1.65706303e-01, 7.83480660e-05, 1.52473276e-03, 3.02816106e-03])

```

```

[107]: extra_attribs = ["rooms_per_hhold", "pop_per_hhold", "bedrooms_per_room"]
#cat_encoder = cat_pipeline.named_steps["cat_encoder"] # old solution
cat_encoder = full_pipeline.named_transformers_["cat"]
cat_one_hot_attribs = list(cat_encoder.categories_[0])
attributes = num_attribs + extra_attribs + cat_one_hot_attribs
sorted(zip(feature_importances, attributes), reverse=True)

```

```

[107]: [(0.3790092248170967, 'median_income'),
(0.16570630316895876, 'INLAND'),
(0.10703132208204354, 'pop_per_hhold'),
(0.06965425227942929, 'longitude'),
(0.0604213840080722, 'latitude'),
(0.054778915018283726, 'rooms_per_hhold'),
(0.048203121338269206, 'bedrooms_per_room'),
(0.04218822024391753, 'housing_median_age'),
(0.015849114744428634, 'population'),
(0.015554529490469328, 'total_bedrooms'),
(0.01524505568840977, 'total_rooms'),
(0.014934655161887776, 'households'),
(0.006792660074259966, '<1H OCEAN'),

```

```
(0.0030281610628962747, 'NEAR OCEAN'),  
(0.0015247327555504937, 'NEAR BAY'),  
(7.834806602687504e-05, 'ISLAND')]
```

```
[108]: final_model = grid_search.best_estimator_  
  
X_test = strat_test_set.drop("median_house_value", axis=1)  
y_test = strat_test_set["median_house_value"].copy()  
  
X_test_prepared = full_pipeline.transform(X_test)  
final_predictions = final_model.predict(X_test_prepared)  
  
final_mse = mean_squared_error(y_test, final_predictions)  
final_rmse = np.sqrt(final_mse)
```

```
[109]: final_rmse
```

```
[109]: 47873.26095812988
```

We can compute a 95% confidence interval for the test RMSE:

```
[110]: from scipy import stats
```

```
[111]: confidence = 0.95  
squared_errors = (final_predictions - y_test) ** 2  
mean = squared_errors.mean()  
m = len(squared_errors)  
  
np.sqrt(stats.t.interval(confidence, m - 1,  
                          loc=np.mean(squared_errors),  
                          scale=stats.sem(squared_errors)))
```

```
[111]: array([45893.36082829, 49774.46796717])
```

We could compute the interval manually like this:

```
[112]: tscore = stats.t.ppf((1 + confidence) / 2, df=m - 1)  
tmargin = tscore * squared_errors.std(ddof=1) / np.sqrt(m)  
np.sqrt(mean - tmargin), np.sqrt(mean + tmargin)
```

```
[112]: (45893.360828285535, 49774.46796717361)
```

Alternatively, we could use a z-scores rather than t-scores:

```
[113]: zscore = stats.norm.ppf((1 + confidence) / 2)  
zmargin = zscore * squared_errors.std(ddof=1) / np.sqrt(m)  
np.sqrt(mean - zmargin), np.sqrt(mean + zmargin)
```

```
[113]: (45893.9540110131, 49773.921030650374)
```

7 Extra material

7.1 A full pipeline with both preparation and prediction

```
[114]: full_pipeline_with_predictor = Pipeline([
        ("preparation", full_pipeline),
        ("linear", LinearRegression())
    ])

full_pipeline_with_predictor.fit(housing, housing_labels)
full_pipeline_with_predictor.predict(some_data)
```

```
[114]: array([ 85657.90192014, 305492.60737488, 152056.46122456, 186095.70946094,
        244550.67966089])
```

7.2 Model persistence using joblib



```
[115]: my_model = full_pipeline_with_predictor
```

```
[116]: #from sklearn.externals import joblib # deprecated, use import joblib instead
import joblib

joblib.dump(my_model, "my_model.pkl") # DIFF
#...
my_model_loaded = joblib.load("my_model.pkl") # DIFF
```

7.3 Example SciPy distributions for RandomizedSearchCV

```
[117]: from scipy.stats import geom, expon
geom_distrib=geom(0.5).rvs(10000, random_state=42)
expon_distrib=expon(scale=1).rvs(10000, random_state=42)
plt.hist(geom_distrib, bins=50)
plt.show()
plt.hist(expon_distrib, bins=50)
plt.show()
```



8 Exercise solutions

8.1 1.

Question: Try a Support Vector Machine regressor (`sklearn.svm.SVR`), with various hyperparameters such as `kernel="linear"` (with various values for the `C` hyperparameter) or `kernel="rbf"` (with various values for the `C` and `gamma` hyperparameters). Don't worry about what these hyperparameters mean for now. How does the best SVR predictor perform?

```
[118]: from sklearn.model_selection import GridSearchCV

param_grid = [
    {'kernel': ['linear'], 'C': [10., 30., 100., 300., 1000., 3000., 10000.
    ↪, 30000.0]},
    {'kernel': ['rbf'], 'C': [1.0, 3.0, 10., 30., 100., 300., 1000.0],
     'gamma': [0.01, 0.03, 0.1, 0.3, 1.0, 3.0]},
    ]

svm_reg = SVR()
grid_search = GridSearchCV(svm_reg, param_grid, cv=5,
    ↪scoring='neg_mean_squared_error', verbose=2, n_jobs=4)
grid_search.fit(housing_prepared, housing_labels)
```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

```
[118]: GridSearchCV(cv=5, estimator=SVR(), n_jobs=4,
    param_grid=[{'C': [10.0, 30.0, 100.0, 300.0, 1000.0, 3000.0,
    10000.0, 30000.0],
    'kernel': ['linear']},
    {'C': [1.0, 3.0, 10.0, 30.0, 100.0, 300.0, 1000.0],
    'gamma': [0.01, 0.03, 0.1, 0.3, 1.0, 3.0],
    'kernel': ['rbf']}],
    scoring='neg_mean_squared_error', verbose=2)
```

The best model achieves the following score (evaluated using 5-fold cross validation):

```
[119]: negative_mse = grid_search.best_score_
rmse = np.sqrt(-negative_mse)
rmse
```

```
[119]: 70286.61835383571
```

That's much worse than the `RandomForestRegressor`. Let's check the best hyperparameters found:

```
[120]: grid_search.best_params_
```

```
[120]: {'C': 30000.0, 'kernel': 'linear'}
```

The linear kernel seems better than the RBF kernel. Notice that the value of `C` is the maximum

tested value. When this happens you definitely want to launch the grid search again with higher values for C (removing the smallest values), because it is likely that higher values of C will be better.

8.2 2.

Question: Try replacing GridSearchCV with RandomizedSearchCV.

```
[121]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import expon, reciprocal

# see https://docs.scipy.org/doc/scipy/reference/stats.html
# for `expon()` and `reciprocal()` documentation and more probability
# → distribution functions.

# Note: gamma is ignored when kernel is "linear"
param_distributions = {
    'kernel': ['linear', 'rbf'],
    'C': reciprocal(20, 200000),
    'gamma': expon(scale=1.0),
}

svm_reg = SVR()
rnd_search = RandomizedSearchCV(svm_reg, param_distributions=param_distributions,
                                n_iter=50, cv=5,
                                → scoring='neg_mean_squared_error',
                                verbose=2, n_jobs=4, random_state=42)
rnd_search.fit(housing_prepared, housing_labels)
```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

```
[121]: RandomizedSearchCV(cv=5, estimator=SVR(), n_iter=50, n_jobs=4,
    param_distributions={'C':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f3be9c1b880>,
    'gamma':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7f3be9b55640>,
    'kernel': ['linear', 'rbf']},
    random_state=42, scoring='neg_mean_squared_error',
    verbose=2)
```

The best model achieves the following score (evaluated using 5-fold cross validation):

```
[122]: negative_mse = rnd_search.best_score_
rmse = np.sqrt(-negative_mse)
rmse
```

```
[122]: 54751.69009488048
```

Now this is much closer to the performance of the RandomForestRegressor (but not quite there yet). Let's check the best hyperparameters found:


```
[123]: rnd_search.best_params_
```

```
[123]: {'C': 157055.10989448498, 'gamma': 0.26497040005002437, 'kernel': 'rbf'}
```

This time the search found a good set of hyperparameters for the RBF kernel. Randomized search tends to find better hyperparameters than grid search in the same amount of time.

Let's look at the exponential distribution we used, with `scale=1.0`. Note that some samples are much larger or smaller than 1.0, but when you look at the log of the distribution, you can see that most values are actually concentrated roughly in the range of $\exp(-2)$ to $\exp(+2)$, which is about 0.1 to 7.4.


```
[124]: expon_distrib = expon(scale=1.)
samples = expon_distrib.rvs(10000, random_state=42)
plt.figure(figsize=(10, 4))
plt.subplot(121)
plt.title("Exponential distribution (scale=1.0)")
plt.hist(samples, bins=50)
plt.subplot(122)
plt.title("Log of this distribution")
plt.hist(np.log(samples), bins=50)
plt.show()
```



The distribution we used for `C` looks quite different: the scale of the samples is picked from a uniform distribution within a given range, which is why the right graph, which represents the log of the samples, looks roughly constant. This distribution is useful when you don't have a clue of what the target scale is:

```
[125]: reciprocal_distrib = reciprocal(20, 200000)
samples = reciprocal_distrib.rvs(10000, random_state=42)
plt.figure(figsize=(10, 4))
plt.subplot(121)
```

```
plt.title("Reciprocal distribution (scale=1.0)")
plt.hist(samples, bins=50)
plt.subplot(122)
plt.title("Log of this distribution")
plt.hist(np.log(samples), bins=50)
plt.show()
```



The reciprocal distribution is useful when you have no idea what the scale of the hyperparameter should be (indeed, as you can see on the figure on the right, all scales are equally likely, within the given range), whereas the exponential distribution is best when you know (more or less) what the scale of the hyperparameter should be.

8.3 3.

Question: Try adding a transformer in the preparation pipeline to select only the most important attributes.

```
[126]: from sklearn.base import BaseEstimator, TransformerMixin

def indices_of_top_k(arr, k):
    return np.sort(np.argpartition(np.array(arr), -k)[-k:])

class TopFeatureSelector(BaseEstimator, TransformerMixin):
    def __init__(self, feature_importances, k):
        self.feature_importances = feature_importances
        self.k = k
    def fit(self, X, y=None):
        self.feature_indices_ = indices_of_top_k(self.feature_importances, self.
→k)
        return self
    def transform(self, X):
```

```
return X[:, self.feature_indices_]
```

Note: this feature selector assumes that you have already computed the feature importances somehow (for example using a `RandomForestRegressor`). You may be tempted to compute them directly in the `TopFeatureSelector`'s `fit()` method, however this would likely slow down grid/randomized search since the feature importances would have to be computed for every hyperparameter combination (unless you implement some sort of cache).

Let's define the number of top features we want to keep:

```
[127]: k = 5
```

Now let's look for the indices of the top k features:

```
[128]: top_k_feature_indices = indices_of_top_k(feature_importances, k)
top_k_feature_indices
```

```
[128]: array([ 0,  1,  7,  9, 12])
```

```
[129]: np.array(attributes)[top_k_feature_indices]
```

```
[129]: array(['longitude', 'latitude', 'median_income', 'pop_per_hhold',
            'INLAND'], dtype='<U18')
```

Let's double check that these are indeed the top k features:

```
[130]: sorted(zip(feature_importances, attributes), reverse=True)[:k]
```

```
[130]: [(0.3790092248170967, 'median_income'),
        (0.16570630316895876, 'INLAND'),
        (0.10703132208204354, 'pop_per_hhold'),
        (0.06965425227942929, 'longitude'),
        (0.0604213840080722, 'latitude')]
```

Looking good... Now let's create a new pipeline that runs the previously defined preparation pipeline, and adds top k feature selection:

```
[131]: preparation_and_feature_selection_pipeline = Pipeline([
        ('preparation', full_pipeline),
        ('feature_selection', TopFeatureSelector(feature_importances, k))
    ])
```

```
[132]: housing_prepared_top_k_features = preparation_and_feature_selection_pipeline.
        ↪fit_transform(housing)
```

Let's look at the features of the first 3 instances:

```
[133]: housing_prepared_top_k_features[0:3]
```

```
[133]: array([[ -0.94135046,  1.34743822, -0.8936472 ,  0.00622264,  1.          ],
              [  1.17178212, -1.19243966,  1.292168   , -0.04081077,  0.          ],
              [  0.26758118, -0.1259716 , -0.52543365, -0.07537122,  1.          ]])
```

Now let's double check that these are indeed the top k features:

```
[134]: housing_prepared[0:3, top_k_feature_indices]
```

```
[134]: array([[ -0.94135046,  1.34743822, -0.8936472 ,  0.00622264,  1.          ],
              [  1.17178212, -1.19243966,  1.292168   , -0.04081077,  0.          ],
              [  0.26758118, -0.1259716 , -0.52543365, -0.07537122,  1.          ]])
```

Works great! :)

8.4 4.

Question: Try creating a single pipeline that does the full data preparation plus the final prediction.

```
[135]: prepare_select_and_predict_pipeline = Pipeline([
        ('preparation', full_pipeline),
        ('feature_selection', TopFeatureSelector(feature_importances, k)),
        ('svm_reg', SVR(**rnd_search.best_params_))
    ])
```

```
[136]: prepare_select_and_predict_pipeline.fit(housing, housing_labels)
```

```
[136]: Pipeline(steps=[('preparation',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('attrs_adder',
                                                                              FunctionTransformer(func=<function add_extra_features at 0x7f3be92d8040>)),
                                                                              ('std_scaler',
                                                                              StandardScaler()))]),
                                                         ['longitude', 'latitude',
                                                                              'housing_median_age',
                                                                              'total_rooms',
                                                                              'total_bedrooms',
                                                                              'population', 'househ...
                                                         TopFeatureSelector(feature_importances=array([6.96542523e-02,
6.04213840e-02, 4.21882202e-02, 1.52450557e-02,
1.55545295e-02, 1.58491147e-02, 1.49346552e-02, 3.79009225e-01,
5.47789150e-02, 1.07031322e-01, 4.82031213e-02, 6.79266007e-03,
1.65706303e-01, 7.83480660e-05, 1.52473276e-03, 3.02816106e-03]),
                                                         k=5)),
                        ('svm_reg',
                        SVR(C=157055.10989448498, gamma=0.26497040005002437))])
```

Let's try the full pipeline on a few instances:

```
[137]: some_data = housing.iloc[:4]
some_labels = housing_labels.iloc[:4]

print("Predictions:\t", prepare_select_and_predict_pipeline.predict(some_data))
print("Labels:\t\t", list(some_labels))
```

```
Predictions:      [ 83384.49158095 299407.90439234  92272.03345144
150173.16199041]
Labels:           [72100.0, 279600.0, 82700.0, 112500.0]
```

Well, the full pipeline seems to work fine. Of course, the predictions are not fantastic: they would be better if we used the best `RandomForestRegressor` that we found earlier, rather than the best `SVR`.

8.5 5.

Question: Automatically explore some preparation options using `GridSearchCV`.

```
[138]: param_grid = [{
    'preparation__num__imputer__strategy': ['mean', 'median', 'most_frequent'],
    'feature_selection__k': list(range(1, len(feature_importances) + 1))
}]

grid_search_prep = GridSearchCV(prepare_select_and_predict_pipeline,
    ↪param_grid, cv=5,
                                scoring='neg_mean_squared_error', verbose=2,
    ↪n_jobs=4)
grid_search_prep.fit(housing, housing_labels)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

```
/usr/local/lib/python3.8/dist-
packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:
9 fits failed out of a total of 240.
The score on these train-test partitions for these parameters will be set to
nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.
```

Below are more details about the failures:

```
-----
9 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.8/dist-
packages/sklearn/model_selection/_validation.py", line 680, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.8/dist-packages/sklearn/pipeline.py", line 390,
in fit
    Xt = self._fit(X, y, **fit_params_steps)
```

```

File "/usr/local/lib/python3.8/dist-packages/sklearn/pipeline.py", line 348,
in _fit
    X, fitted_transformer = fit_transform_one_cached(
File "/usr/local/lib/python3.8/dist-packages/joblib/memory.py", line 349, in
__call__
    return self.func(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/sklearn/pipeline.py", line 893,
in _fit_transform_one
    res = transformer.fit_transform(X, y, **fit_params)
File "/usr/local/lib/python3.8/dist-packages/sklearn/base.py", line 855, in
fit_transform
    return self.fit(X, y, **fit_params).transform(X)
File "<ipython-input-126-6a801ecaa128>", line 14, in transform
IndexError: index 15 is out of bounds for axis 1 with size 15

```

```

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_search.py:969:
UserWarning: One or more of the test scores are non-finite: [nan nan nan nan nan
nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan
nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan
nan nan nan nan nan nan nan nan nan nan nan nan nan nan]
warnings.warn(

```

```

[138]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('preparation',
                                             ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='median')),
                  ('attrs_adder',
                     FunctionTransformer(func=<function add_extra_features at
0x7f3be92d8040>)),
                  ('std_scaler',
                     StandardScaler()))]),
                  ['longitude',
                  'latitude',
                  'housing_median_age',
                  'total_rooms',
                  'total_be...
                    5.47789150e-02, 1.07031322e-01, 4.82031213e-02, 6.79266007e-03,
                    1.65706303e-01, 7.83480660e-05, 1.52473276e-03, 3.02816106e-03]),
                                                              k=5)),
                  ('svm_reg',
                   SVR(C=157055.10989448498,
                      gamma=0.26497040005002437))]),
                  n_jobs=4,
                  param_grid=[{'feature_selection__k': [1, 2, 3, 4, 5, 6, 7, 8, 9,
                                                            10, 11, 12, 13, 14, 15, 16],

```

```
        'preparation__num__imputer__strategy': ['mean',
                                                  'median',
        'most_frequent']]],
        scoring='neg_mean_squared_error', verbose=2)
```

```
[139]: grid_search_prep.best_params_
```

```
[139]: {'feature_selection__k': 1, 'preparation__num__imputer__strategy': 'mean'}
```

The best imputer strategy is `most_frequent` and apparently almost all features are useful (15 out of 16). The last one (ISLAND) seems to just add some noise.

Congratulations! You already know quite a lot about Machine Learning. :)